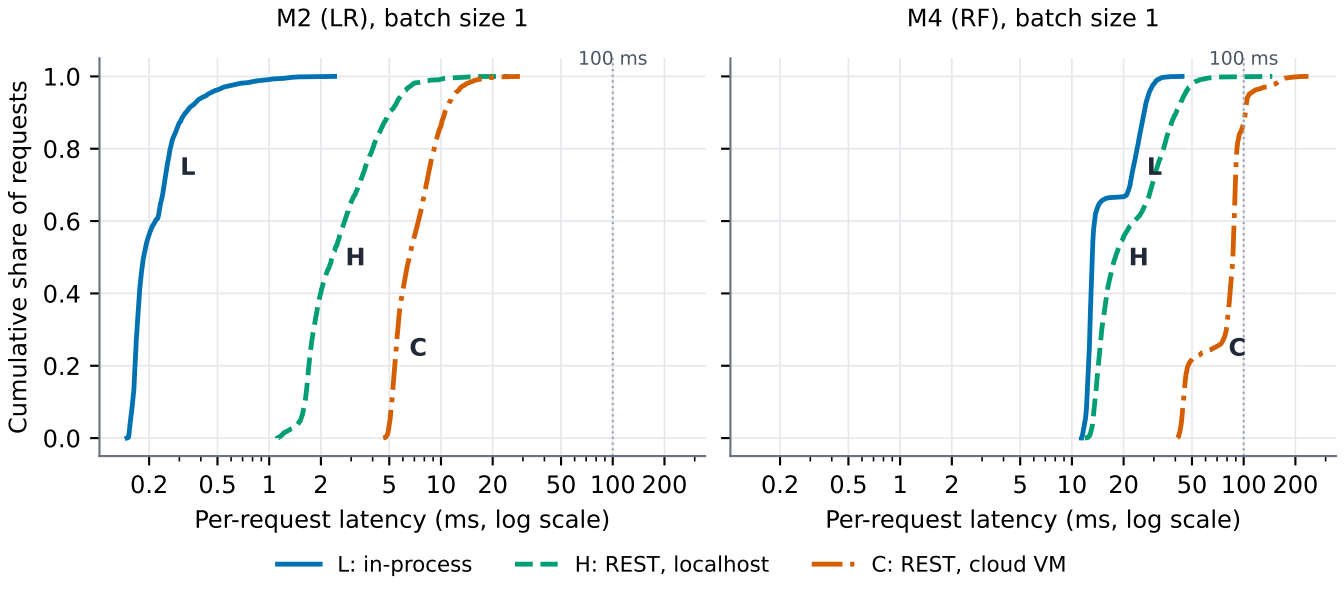
M2 (LR), batch size 1
M4 (RF), batch size 1
100 ms
100 ms
Cumulative share of requests
1.0
0.8
0.6
0.4
0.2
0.0
L
H
C
L
H
C
0.2
0.5
1
2
5
10
20
50
100
200
Per-request latency (ms, log scale)
Per-request latency (ms, log scale)
L: in-process
H: REST, localhost
C: REST, cloud VM